

Joseph Okonda L.
Stanford University

1

2. Related Work

Object Detection systems fall into two main recent categories: Region Proposal based detectors (R-CNN) [3, 4, 14] detect objects by first running a proposal generator which outputs regions of the image that are likely to contain an object of interest. The proposals are then further processed to classify any objects that they contain and to refine the objects’ bounding boxes. Such systems give the best accuracy but also have a high latency. Single Shot Detectors (SSDs) [12, 9] on the other hand frame object detection as a regression problem. By directly predicting the coordinates of the bounding boxes, SSDs are able to detect objects without extracting the preliminary region proposal as in R-CNNs. This makes them quite fast. However, their low latency comes at the cost of reduced accuracy.

In this project we use a single shot detector to spot nutrition information. We make this choice because we’d like our models to be able to perform inference in real time.

Machine Reading systems range from traditional OCR systems such as [1] to recent models capable of recognizing text in unconstrained scene images like [7, 6]. Traditional systems impose strict constraints on their inputs. In particular, they assume that input images are printed documents on a clear background without much noise. This makes them unfavorable for the task at hand because even though the texts we’d like to read are machine generated, they occur in very noisy environments with auxiliary text that we’d like to ignore. Furthermore, the data we’d like to extract has structure that we’d like to preserve. Unconstrained text spotting and classification is a much harder problem than simple optical character recognition. Unlike OCR, this task imposes no constraints on the input image [6]. Models that solve this problem usually cast it as an object detection task where various approaches are used to localize and classify the characters present in the input image. Such models offer much versatility in terms of their inputs but have high latency especially when the input images are rich in text.

The task in this project is much more complex than traditional OCR but simpler than unconstrained text spotting. We cast the transcription task as an image captioning task where the ‘caption’ to be generated is the text contained in the image.

3. Methods

3.1. Data

We begin by collecting data in the form of short (15 - 20 sec.) videos shot from a cellphone camera¹. We collect 700 such videos. We then turn the videos into images using FFmpeg ensuring that we only sample a few images (5) every second. This is done to ensure that images from

the same video have as much variance as possible. After this operation, every video is reduced to approximately 80 images giving us a dataset of ≈ 40000 images. Each image is labelled with ground truth bounding polygons as shown in Figure 2. Additionally, we transcribe every *video*’s list of ingredients and its nutrition facts dictionary. We split our data into train-dev-test sets using a 95-2.5-2.5 scheme. Crucially, this split is done at the video level instead of the image level. This is to ensure that our test and development sets are truly unseen during training.

3.2. Spotting

For spotting, we repurpose the YOLO9000 model [12] to fit our task. In particular, we only adopt the first 13 layers. We add on a final convolution layer with as many filters as the size of our target tensor. The receptive field of the kernels in this layer is the entire input filtermap. We do this in order to avoid including fully connected layers in the architecture. After every convolution layer we add Batch Normalization [5] in order to speed up training. We also include dropout [16] after every batch normalization module for regularization. Table 1 shows our architecture. We initialize the models with YOLO9000 weights trained on the MS-COCO [8] dataset. We freeze the first 7 layers of the network. Our goal is train this network to predict a tensor of dimension $S \times S \times 7 \times B$ where in our case $S = 3$, i.e the number of grid cells along a single dimension of the input image and $B = 2$ i.e each grid cell predicts 2 bounding boxes. The 7 in the output tensor comes from the fact that we predict an ‘object score’ telling us if we have an object of interest in a given grid cell, 2 center coordinates, a length, a width, and the probability of the object in that grid cell belonging to one of 2 classes. We train the network to minimize YOLO’s multipart MSE loss [12] using the Adam optimizer with a learning rate of 1.5×10^{-6} . To evaluate our model’s performance we measure the mean average precision of our bounding box predictions with 0.7 intersection over union (IoU) as our correctness threshold.

3.3. Transcribing

For transcription, we begin by extracting the regions of interest. We do this by using the ground truth bounding boxes of our data to crop out the list of ingredients and the nutrition fact-box. This yields two new images per image, each of which is resized to be 256×256 and then fed *separately* into our encoder. Figure 3 shows example input images to our transcription module. Our encoder is a 14-layer convolutional network whose first 10 layers are initialized from a pretrained YOLO model. However, unlike in the localization module, we do not freeze any layers in the network. Our decoder is a 3-layer GRU model [2] with a hidden state of size 1024. We begin by pre-training the encoder to solve the binary classification task of differ-

¹<https://youtu.be/VssRmq-MHBc> Here is a sample video


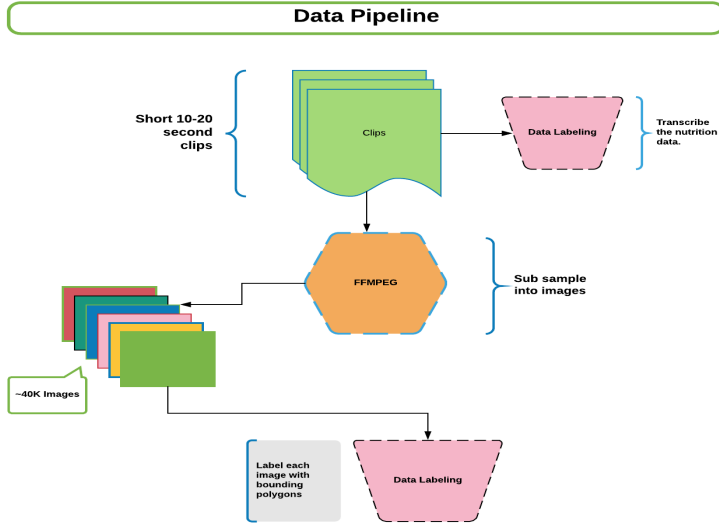


Figure 2. **Left:** Our Data Collection Pipeline. **Right:** An example image from our dataset with Bounding Polygons in blue and YOLO's grid cells in red.

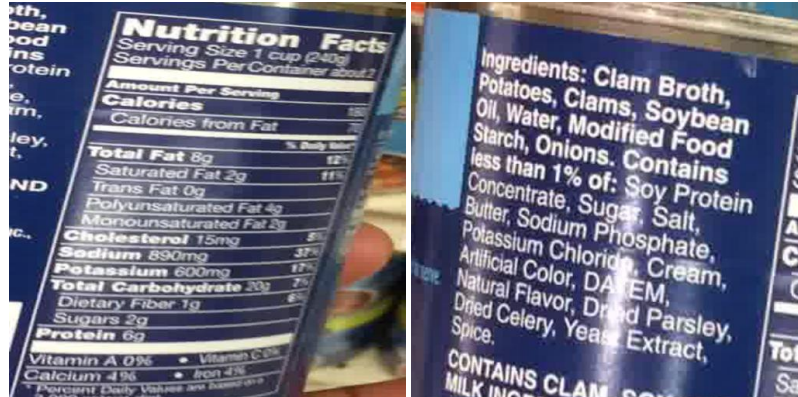


Figure 3. The input images to the Transcription Network. We use the ground truth bounding boxes to extract the regions of interest and resize them to be 256×256 . **Top:** A closely cropped nutrition facts section. The **Bottom:** A closely cropped list of ingredients. They are 2 different images.

entiating between a list of ingredients and a nutrition fact-box. Without this pre-training step, end-to-end optimization quickly becomes unstable. We train the network to minimize per-word cross entropy loss using Adam with a learning rate of 1.5×10^{-3} .

$$\ell = - \sum_{i=0}^m \sum_{t=0}^T y'^{<t>} \log(y^{<t>})$$

During training we employ ‘teacher forcing’ [15] where at time t , we feed in the ground truth word vector of time $t-1$. Therefore, the transcription module is attempting to model the following conditional probability:

$$\log p(Y|F(X)) = \sum_{t=0}^T \log p(y^{<t>}|F(X), y^{<t-1>})$$

Type	Filters	size or stride
Conv	32	3×3
MaxPool		$2 \times 2/s = 2$
Conv	64	3×3
MaxPool		$2 \times 2/s = 2$
Conv	128	3×3
Conv	64	3×3
Conv	128	3×3
MaxPool		$2 \times 2/s = 2$
Conv	256	3×3
Conv	128	3×3
Conv	256	3×3
MaxPool		$2 \times 2/s = 2$
Conv	512	3×3
Conv	256	3×3
Conv	512	3×3
Conv	256	3×3
Conv	512	3×3
MaxPool		$2 \times 2/s = 2$
Conv	$S*S*B*C$	$K \times K$

Table 1. Architecture of the Localization Network

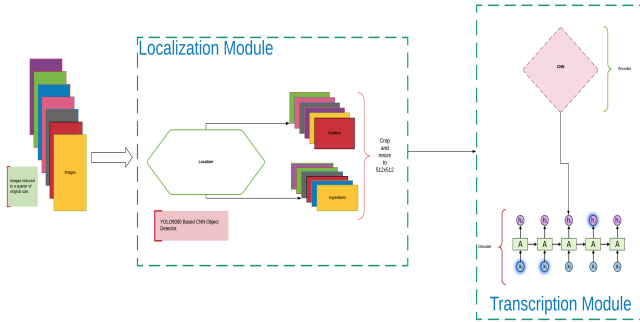


Figure 4. A cartoon of our Entire pipeline.

Where F is the encoder and X is the raw input image. To evaluate our module, we measure the BLEU score [11] using samples from our model and the ground truth annotations. Figure 4 shows an image of our full pipeline.

4. Results

4.1. Spotting Results.

Our trained detection module successfully learns to predict the bounding boxes. Figure 5 shows the results from training the detection network for 100 epochs. When cal-

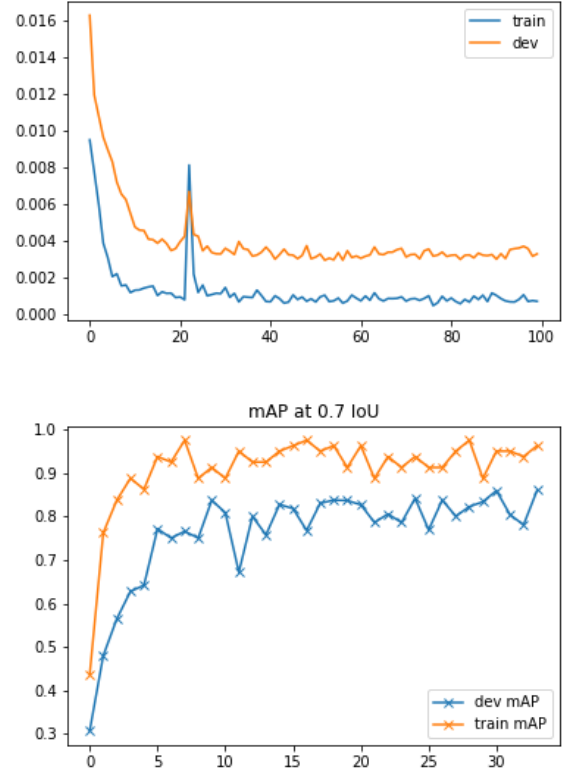


Figure 5. **Left:** Loss curves from both our training and development sets. **Right:** Mean Average Precision Curve measured at 0.7 IoU

culating the mAP scores, we set the ‘objectness threshold’ for each grid cell to 0.9 and use 0.7 as the threshold for each class. We qualitatively evaluate our model by visualizing some of the predicted bounding boxes. Figure 6 shows some of these qualitative results.

4.2. Transcription Results.

The results of our transcription module are not impressive. Figure 8 shows the BLEU scores on both the training and development set. While our model learns to generate text that is consistent with our target grammar, as shown in figure 7, it fails to correctly condition on the input image. A symptom of this is that the transcriptions the model struggles when transcribing numeric values. These are values that the model cannot transcribe by solely relying on a simple language model.

5. Conclusion & Future Work

So far, we have successfully trained a YOLO based network capable of detecting patches of an input image that contain either the list of ingredients or the nutrition fact-box. Additionally, we train an Encoder-Decoder model that




Figure 6. Example bounding boxes generated by our model. The red boxes are the ground truth bounding boxes while the blue boxes are predicted by our detector.

```

Truth
---
<start> {
  num_servings : 2 8 serving_size : 4 0 g
  calories_per_serving : 1 6 0
  calories_from_fat : 4 0
  total_fat : 4 . 5 g sat_fat : 1 . 5 g
  trans_fat : 0 g poly_fat : 1 g
  mono_fat : 2 g cholesterol : 0 mg
  sodium : 4 1 0 mg potassium : 6 0 mg
  total_carb : 2 5 g dietary_fiber : 1 g
  total_sugar : 1 g added_sugar : n/a
  protein : 3 g
} <end>
Prediction
---
<start> {
  num_servings : 4 serving_size : 1 4 0 g
  calories_per_serving : 1 4 0
  calories_from_fat : 4 5
  total_fat : 4 . 5 g sat_fat : 0 . 5 g
  trans_fat : 0 g poly_fat : n/a
  mono_fat : n/a
  cholesterol : 0 mg
  sodium : 1 0 0 mg potassium : n/a
  total_carb : 2 4 g dietary_fiber : 2 g
  total_sugar : 1 3 g added_sugar : n/a
  protein : 1 4 g
} <end>

```

Figure 7. Example Transcription Results **Top**: Ground truth transcription, **Bottom**: Sampled transcriptions.




Figure 8. Train and Dev BLEU scores

attempts to transcribe the nutrition information present in a closely cropped image patch. We are less successful at the second task.

From our analysis, we discovered that the transcription network has a hard time conditioning its on the input image’s encodings. One way to solve this problem would be by incorporating an attention mechanism [10] in the decoder. Moreover, at the moment, due to the way we collect and annotate our transcription data (We annotate each video), we effectively end up with a very small vocabulary. This means that even if the transcription model improves significantly, we’ll be overfitting to a very small vocabulary. Our transcription module also struggles when attempting to transcribe blurry input images. The transcription task requires the model to learn fine-grained details in the input, therefore having blurry input makes this much harder.

Unfortunately, most images sampled from videos have motion blur. This could be remedied by having an explicit pre-processing module that does the deblurring.

Finally, the main bottleneck of our project is the Localizer module. This task requires us to label bounding boxes for a large number of images. Therefore, a future work will focus on ways of replacing this module with a different model that does not require bounding box annotations.

Acknowledgments

We acknowledge Google Cloud and AWS for providing compute infrastructure and the entire 231N teaching team for organizing the class.

References

- [1] T. M. Breuel. The ocrpus open source ocr system. In *DRR*, 2008. 2
- [2] K. Cho, B. van Merriënboer, aglar Gülehre, D. Bahdanau, F. Bougares, H. Schwenk, and Y. Bengio. Learning phrase representations using rnn encoder-decoder for statistical machine translation. In *EMNLP*, 2014. 2
- [3] R. B. Girshick. Fast r-cnn. *2015 IEEE International Conference on Computer Vision (ICCV)*, pages 1440–1448, 2015. 2
- [4] K. He, G. Gkioxari, P. Dollár, and R. B. Girshick. Mask r-cnn. *2017 IEEE International Conference on Computer Vision (ICCV)*, pages 2980–2988, 2017. 2
- [5] S. Ioffe and C. Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. In *ICML*, 2015. 2
- [6] M. Jaderberg, K. Simonyan, A. Vedaldi, and A. Zisserman. Reading text in the wild with convolutional neural networks. *International Journal of Computer Vision*, 116:1–20, 2015. 2
- [7] M. Jaderberg, A. Vedaldi, and A. Zisserman. Deep features for text spotting. In *ECCV*, 2014. 2
- [8] T.-Y. Lin, M. Maire, S. J. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick. Microsoft coco: Common objects in context. In *ECCV*, 2014. 2
- [9] W. Liu, D. Anguelov, D. Erhan, C. Szegedy, S. E. Reed, C.-Y. Fu, and A. C. Berg. Ssd: Single shot multibox detector. In *ECCV*, 2016. 2
- [10] T. Luong, H. Pham, and C. D. Manning. Effective approaches to attention-based neural machine translation. In *EMNLP*, 2015. 5
- [11] K. Papineni, S. Roukos, T. Ward, and W.-J. Zhu. Bleu: a method for automatic evaluation of machine translation. In *ACL*, 2002. 4
- [12] J. Redmon, S. K. Divvala, R. B. Girshick, and A. Farhadi. You only look once: Unified, real-time object detection. *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 779–788, 2016. 2
- [13] J. Redmon and A. Farhadi. Yolo9000: Better, faster, stronger. *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 6517–6525, 2017. 1
- [14] S. Ren, K. He, R. B. Girshick, and J. Sun. Faster r-cnn: Towards real-time object detection with region proposal networks. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 39:1137–1149, 2015. 2
- [15] J. Schmidhuber. Deep learning in neural networks: An overview. *Neural networks : the official journal of the International Neural Network Society*, 61:85–117, 2015. 3
- [16] N. Srivastava, G. E. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov. Dropout: a simple way to prevent neural networks from overfitting. *Journal of Machine Learning Research*, 15:1929–1958, 2014. 2
- [17] I. Sutskever, O. Vinyals, and Q. V. Le. Sequence to sequence learning with neural networks. In *NIPS*, 2014. 1